

# Module 12 Pre-Read - Docker Part 2

Wednesday, June 5, 2024 05:55

## Pushing an image to DockerHub

Make sure the image has a correct tag (word that describes the image)

```
docker build -t decathinator/image_name:latest
```

Format is docker hub username, slash, the image name, and then optional description word

```
docker push decathinator/image_name
```

Push to dockerhub

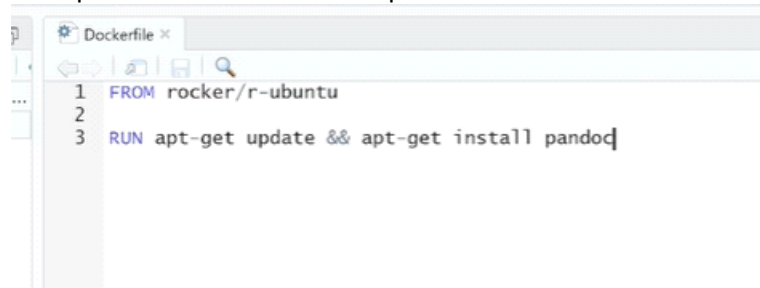
## R in Docker, the easy way

There is a group called Rocker on the DockerHub website that have very useful images

```
docker pull rocker/r-ubuntu
```

In console. Now we have installed the rocker/r-ubuntu image

Example of docker file after we pulled that shows how we can build on the image:



## RStudio IDE in Docker

```
docker pull rocker/rstudio
```

Pull the docker image

```
docker run -e PASSWORD="secret123" -p 8787:8787 rocker/rstudio
```

-e sets the environment variable password

-p is port option

(optional) -d means detached mode, means that container continues running in the background

Navigate to <http://localhost:8787>

Username is rstudio

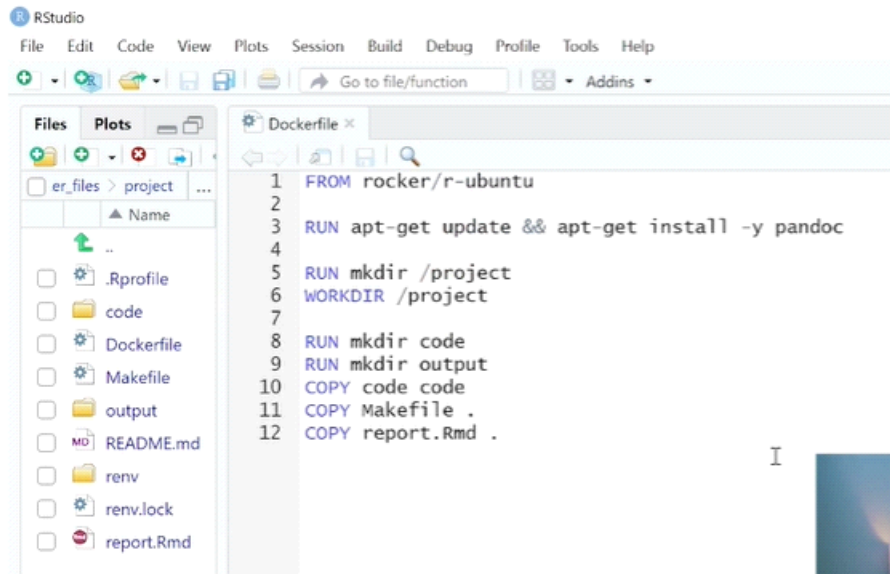
Password is secret123

To shut down the container, you can just close the terminal window or type `docker kill mycontainerid` (you can find the container id if you do `docker container ls`)

## Building a fully automated workflow with R and docker

Navigate to project directory and make a dockerfile via `touch dockerfile`

Example of a dockerfile:



1. Navigate to project directory
2. `docker build -t project_image .`
3. `docker run -it project_image bash`  
 Check contents to see if contents were copied correctly  
 Now you can run `bash` and check via `ls`

### renv and Docker

Need to manually copy over the renv files except for all the stuff in renv library

In the dockerfile:

```
13
14 COPY .Rprofile .
15 COPY renv.lock .
16 RUN mkdir renv
17 COPY renv/activate.R renv
18 COPY renv/settings.dcf renv
19
20 RUN Rscript -e "renv::restore(prompt = FALSE)"
```

### Mounting directories

I want the directory on container to sync with directory on local computer

`docker run -it -v <path to local directory>:<path to image directory>`

The `-v` option mounts the directory

On a windows machine, I need to do `docker run -it -v "//c/<path to local directory>":<path to image directory>`

After you build the image, you should be able to view whatever file

### Using command substitution to shorten docker run command

Instead of writing out the whole file path, you can use a command substitution

`docker run -it -v $(pwd):<path to image directory> project_image bash`

For windows, need to do `docker run -it -v "$($($pwd))/final_report:<path to image directory/final_report> project_image bash`

### Finalizing the automated build

Need to create an entry point to the container

In the dockerfile:

```
23
24 CMD make && mv report.html final_report
25
```

Now the report should be made automatically

## A Make rule for docker builds

In the makefile:

```
11
12 # DOCKER-ASSOCIATED RULES
13 PROJECTFILES = report.Rmd code/01_make_output.R code/02_render_report.R Makefile
14 RENVFILES = renv.lock renv/activate.R renv/settings.dcf
15
16 # rule to build image
17 - project_image: Dockerfile $(PROJECTFILES) $(RENVFILES)
18   docker build -t project_image .
19   touch $$@
20
```

\$\$ evaluates to the target name in make

This way if any stuff changes then the image will rebuild

## A Make rule for docker run

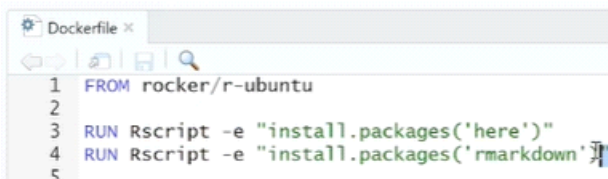
In the makefile:

```
20
21 # rule to build the report automatically in our container
22 - final_report/report.html: project_image
23   docker run -v "$$(pwd)/final_report":/project/final_report project_image
24
25
```

Make requires two dollar signs not just one

## An alternative to renv

Do manual package installation inside the dockerfile



```
Dockerfile x
1 FROM rocker/r-ubuntu
2
3 RUN Rscript -e "install.packages('here')"
4 RUN Rscript -e "install.packages('rmarkdown')"
```

## Contrasting renv and install.packages for use with Docker

### renv

#### PRO

- renv.lock gives record of package versions
- Package versions are same across multiple builds
- Easy to synchronize docker images with local computer

#### CON

- Docker builds don't cache efficiently
- Careful Dockerfile coding required

---

### install.packages()

#### PRO

- Docker builds efficiently use cache
- No need to understand renv intricacies

#### CON

- Package versions may change across builds
- (Slightly) harder to list of package versions
- Harder to synchronize versions across local computer and Docker image



## Create an ad-hoc lock file

```

Terminal 1 - MINGW64/c/Users/David Benkeser/Documents/docker_files/project
$ docker run project_image Rscript -e "installed.packages()[, 'version']"
base64enc    bslib      cachem      digest     evaluate   fastmap     fs
"0.1-3"      "0.4.1"      "1.0.6"     "0.6.30"    "0.18"     "1.1.0"     "1.5.2"
glue         here         highr      htmltools   jquerylib   jsonlite    knitr
"1.6.2"      "1.0.1"      "0.9"       "0.5.3"     "0.1.4"    "1.8.3"     "1.40"

$ docker run project_image Rscript -e "installed.packages()[, 'version']" > bootleg.lock

```

So this is when you aren't using renv, and showing that we can still retrieve this information

### A containerized development workflow using the RStudio IDE

Mount everything and edit and work interactively etc

```

Console  Terminal x Background Jobs x
Terminal 1 - MINGW64/c/Users/David Benkeser/Documents/docker_files/project
$ docker run -e PASSWORD="secret123" -p 8787:8787 -v "/$(pwd)":/home/rstudio rocker/rstudio

```

And once you save the file in the container, it will also update locally